

Implement K-Nearest Neighbors algorithm on diabetes.csv dataset.

Compute confusion matrix, accuracy, error rate, precision and recall on the given dataset.

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, ConfusionMatrixDisplay, classification_report
import matplotlib.pyplot as plt
```

Dataset Description:

We will try to build a machine learning model to accurately predict whether or not the patients in the dataset have diabetes or not? The datasets consists of several medical predictor variables and one target variable, Outcome. Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and so on.

```
In [3]: df = pd.read_csv('diabetes.csv')
df.head()
```

```
Out[3]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
0	6	148	72	35	0	33.6	0.627	51	1
1	1	85	66	29	0	26.6	0.351	33	0
2	8	183	64	0	0	23.3	0.672	33	1
3	1	89	66	23	94	28.1	0.167	34	1
4	0	137	40	35	168	43.1	2.288	33	1

```
In [4]: df = pd.read_csv('diabetes.csv')
df.head()
```


```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
0	6	148	72	35	0	33.6	0.627	51	1
1	1	85	66	29	0	26.6	0.351	33	0
2	8	183	64	0	0	23.3	0.672	33	1
3	1	89	66	23	94	28.1	0.167	34	1

4 0 137 40 35 168 43.1 2.288 :

```
In [5]: x = df.drop('Outcome', axis=1)
        y = df['Outcome']
        sns.countplot(x=y)
```

Out[5]: <Axes: xlabel='Outcome', ylabel='count'>



```
In [6]: y.value_counts()
```

Out[6]: 0 500
 1 268
 Name: Outcome, dtype: int64

```
In [7]: scaler = MinMaxScaler()
        x_scaled = scaler.fit_transform(x)
        x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size=
        x.shape
```

Out[7]: (768, 8)

```
In [8]: print("x_train.shape : ", x_train.shape, "\nx_test.shape : ", x_test.shape)
```

x_train.shape : (537, 8)
x_test.shape : (231, 8)

```
In [9]: knn = KNeighborsClassifier(n_neighbors = 5)
        knn.fit(x_train, y_train)
```


```
Out[9]: KNeighborsClassifier()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [10]: y_pred = knn.predict(x_test)
ConfusionMatrixDisplay.from_predictions(y_test,y_pred)
```

```
Out[10]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7d331401bdf0>
```



```
In [11]: print(classification_report(y_test,y_pred))
```


	precision	recall	f1-score	support
0	0.74	0.78	0.76	151
1	0.54	0.49	0.51	80
accuracy			0.68	231
macro avg	0.64	0.63	0.64	231
weighted avg	0.67	0.68	0.68	231

```
In [12]: error = []
for k in range(1,41):
    knn = KNeighborsClassifier(n_neighbors = k)
    knn.fit(x_train,y_train)
    pred = knn.predict(x_test)
    error.append(np.mean(pred != y_test))
```

```
In [13]: plt.figure(figsize = (16,8))
plt.xlabel('Value of K')
```

```
plt.ylabel('Error')
plt.grid()
plt.xticks(range(1,41))
plt.plot(range(1,41),error,marker = '.')
```

Out[13]: [



In [14]: `knn=KNeighborsClassifier(n_neighbors = 33)`
`knn.fit(x_train,y_train)`

Out[14]: `KNeighborsClassifier(n_neighbors=33)`

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [15]: `y_pred = knn.predict(x_test)`
`print(classification_report(y_test,y_pred))`

	precision	recall	f1-score	support
0	0.77	0.87	0.81	151
1	0.67	0.50	0.57	80
accuracy			0.74	231
macro avg	0.72	0.68	0.69	231
weighted avg	0.73	0.74	0.73	231